

# Have a Seat on the ErasureBench: Easy Evaluation of Erasure Coding Libraries for Distributed Storage Systems

2016 Workshop on Planetary-Scale Distributed Systems

**Sébastien Vaucher**, Hugues Mercier, Valerio Schiavoni

Institute of Computer Science  
Université de Neuchâtel, Switzerland

sebastien.vaucher@unine.ch

Budapest, Hungary, 26 September 2016



# Motivation

More and more data needs to be stored reliably on online servers.  
Reliability can be provided through:

- Replication
- Erasure coding

# Motivation


The characteristics of an erasure coding algorithm are difficult to evaluate.

Evaluation is often done theoretically or by simulation.

# Erasure coding

Goal: add redundancy to cope with data loss/corruption

Example using a  $(5, 2)$  Reed-Solomon code:



# Key features

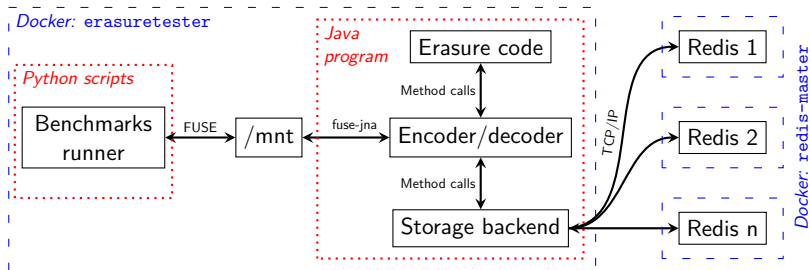
- Compatible with existing benchmark programs
- Automated benchmarks execution
- Containerized storage nodes ( $> 1$  per physical node)
- Replay fault traces

# Evaluation example

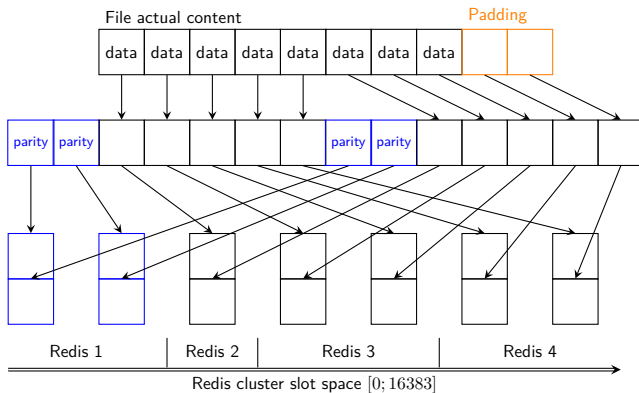
How to evaluate a new algorithm.

1. Program the algorithm as a Java class
2. Write benchmarks as Python functions
  - Debian-compatible programs can be launched as sub-processes
3. Configure the evaluation
  - e.g. algorithm parameters, fault trace, ...
4. Easily deploy the solution to a Docker cluster
5. Collect results

# Technical components



# Blocks distribution





# Metadata management

Each block is identified by a 32-bit key. Using it, we derive:

1. Key of the blocks aggregation stored in Redis
2. Offset within that aggregation

The list of all block keys is kept in memory.

# Automated deployment and scaling

As part of ERASUREBENCH, we provide scripts that automate the deployment of the solution to a Docker Swarm cluster, up to the collection of results.

# Evaluation

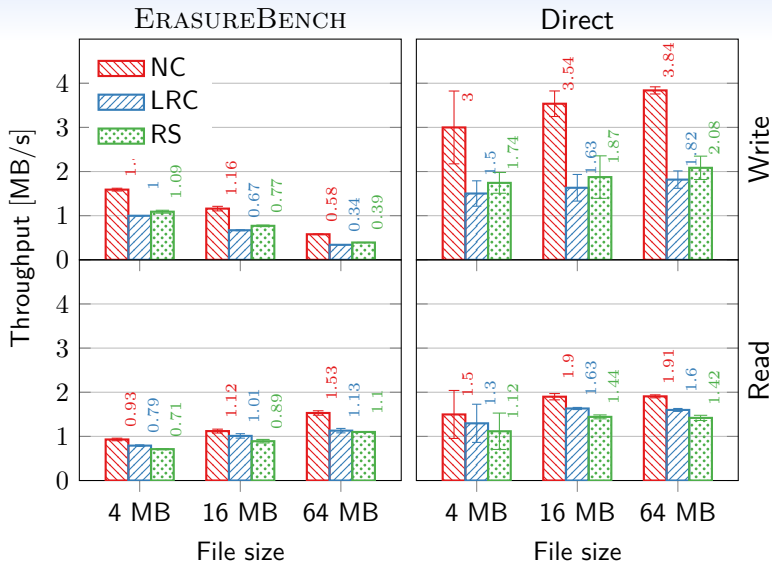
We evaluated algorithms from "XORing Elephants: Novel Erasure Codes for Big Data".

NC No erasure coding

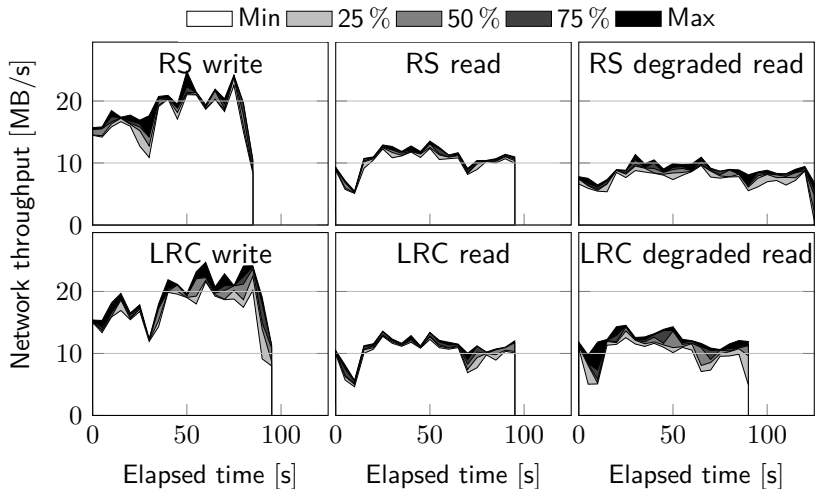
RS Reed-Solomon (10, 4)

LRC Locally Repairable Code (10, 6, 5)

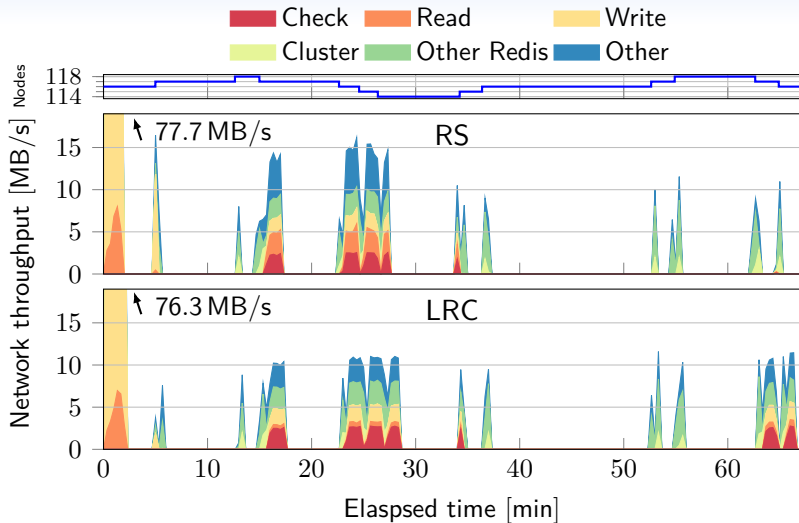
# Throughput



# Traffic



# Trace



# Conclusion

Using ERASUREBENCH, evaluating an erasure coding algorithm under real conditions is easier and cheaper.

Available open-source at  
<https://github.com/safecloud-project/erasurebench>