

Have a Seat on the ErasureBench: Easy Evaluation of Erasure Coding Libraries for Distributed Storage Systems

2016 Workshop on Planetary-Scale Distributed Systems

Sébastien Vaucher, Hugues Mercier, Valerio Schiavoni

Institute of Computer Science
Université de Neuchâtel, Switzerland

sebastien.vaucher@unine.ch

Budapest, Hungary, 26 September 2016



Outline

- Motivation
- Recap about erasure coding
- ERASUREBENCH
- Evaluation
- Conclusion

Motivation

More and more data needs to be stored reliably on online servers.
Reliability can be provided through:

- Replication
- Erasure coding


Motivation

- The characteristics of erasure coding algorithms are difficult to evaluate (encoding, decoding, complexity, latency, ...)
- Evaluation is often done theoretically or by simulation

Erasure coding

Goal: add redundancy to cope with data loss/corruption

Example using a $(5, 2)$ Reed-Solomon code:



ERASUREBENCH key features

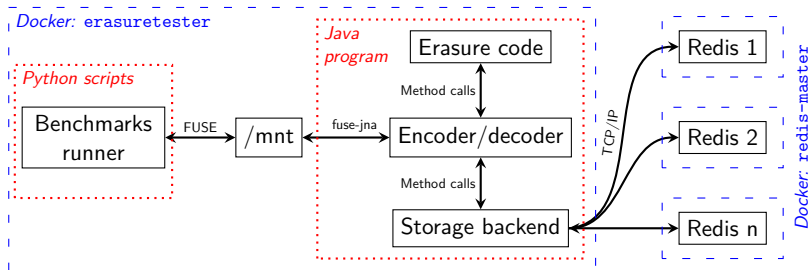
- Compatible with existing benchmark programs
- Automated benchmarks execution
- Containerized storage nodes (> 1 per physical node)
- Can replay fault traces

Evaluation example

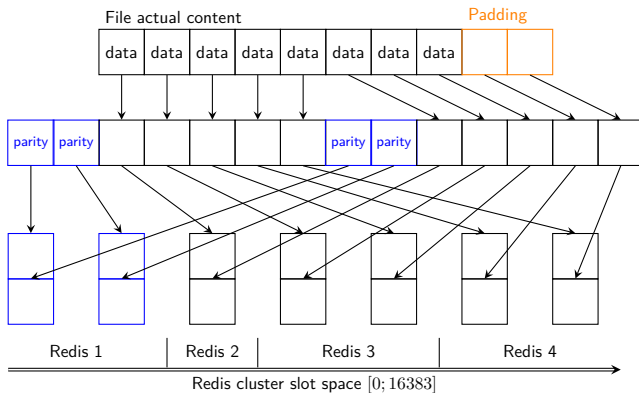
How to evaluate a new erasure coding algorithm

1. Program the algorithm as a Java class
2. Write benchmarks as Python functions
 - Debian-compatible programs can be launched as sub-processes
3. Configure the evaluation
 - e.g. algorithm parameters, fault trace, ...
4. Easily deploy the solution to a Docker cluster
5. Collect results

ERASUREBENCH technical components



Blocks distribution



ERASUREBENCH metadata management

- Each block is identified by a 32-bit key. Using it, we derive:
 1. Key of the blocks aggregation stored in Redis
 2. Offset within that aggregation
- The list of all block keys is kept in memory

ERASUREBENCH automated deployment and scaling

As part of ERASUREBENCH, we provide scripts that automate the deployment of the solution to a Docker Swarm cluster, up to the collection of results

Evaluation

We evaluated algorithms from the following paper:

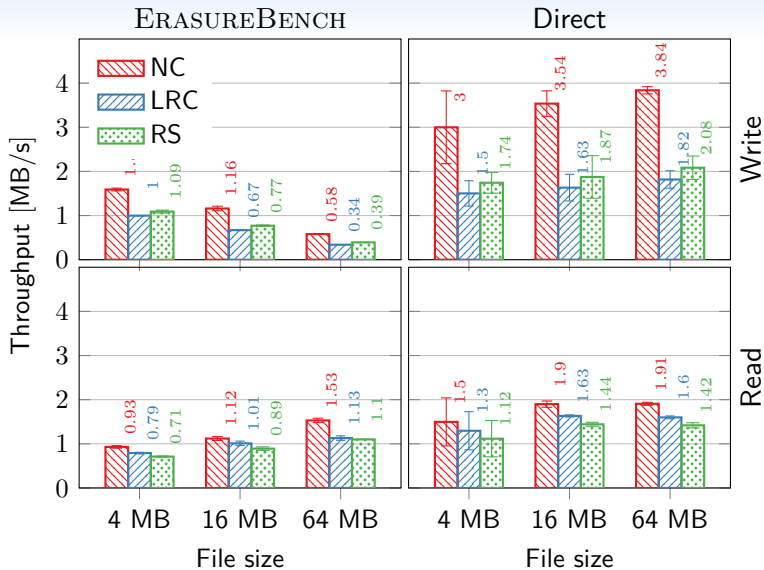
M. Sathiamoorthy, M. Asteris, D. Papailiopoulos, A. G. Dimakis, R. Vadali, S. Chen, and D. Borthakur, "XORing elephants: Novel erasure codes for big data," in Proceedings of the VLDB Endowment, vol. 6, 2013, pp. 325–336.

NC No erasure coding

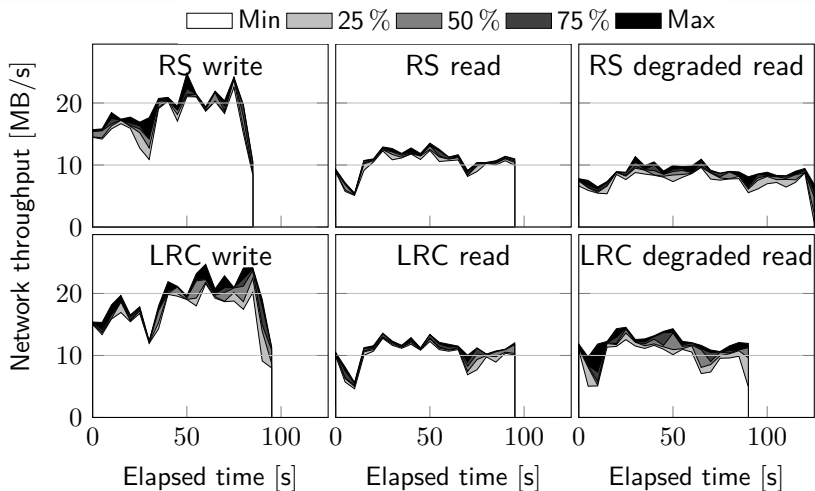
RS $(10, 4)$ Reed-Solomon code

LRC $(10, 6, 5)$ Locally Repairable Code

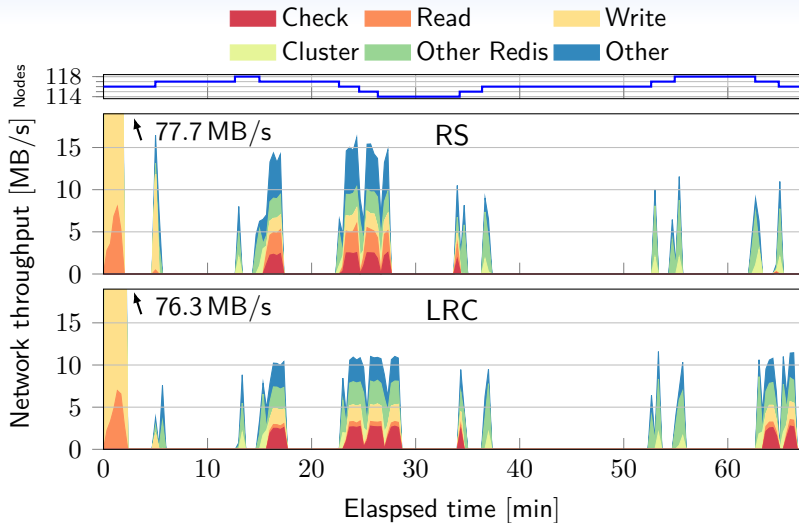
Throughput



Traffic



Trace



Limitations

- Algorithms need to be written in Java
- High memory consumption
- Strong dependency on Docker and Redis

Conclusion

Using ERASUREBENCH, evaluating an erasure coding algorithm under real conditions is easier and cheaper

Available open-source at
<https://github.com/safecloud-project/erasurebench>